Speaker-Embedding and Change-Score Sensitivity at Nepali–English Code-Switch Points

Ishan Gautam
Pulchowk Campus, Institute of Engineering
Pulchowk, Lalitpur
Tribhuvan University, Nepal
isangautam@gmail.com

***Abstract*—People who code-switch change language mid-utterance without changing who they are. We ask whether a Nepali–English switch within one speaker lowers speaker-embedding similarity and raises a speaker-change score. In a pre-specified analysis of real lecture speech from 17 YouTube channels (a proxy for speakers), we compare same-speaker windows that straddle a switch with matched same-language controls, for three speaker embeddings. Between language runs of at least 1 s, similarity is lower across a switch for all three embeddings and in every channel (d_z 0.19–0.27), and cross-language same-speaker scores are 0.17–0.22 SD lower. A two-window change score is slightly higher at these switches (AUC 0.57–0.58, in a primary cell chosen after a pilot), but over all switches, most of which border shorter runs, there is no positive ranking separation (pooled AUC ≈ 0.49). For two embeddings the excess survives the planned adjustment for measured acoustic covariates, although the observational design cannot show that language causes it. On Hindi– and Bengali–English (MUCS 2021) the direction is the same. Two mitigations meant to remove language information reduced the excess by 18.9% and 24.9%, inconclusively, and one made language more, not less, decodable. In a post hoc check, two off-the-shelf diarization pipelines almost never reported a false speaker change at natural switches. The similarity drop is consistent but small, the evidence is preliminary (no independent listener has checked the switch points), and whether it matters in conversational diarization remains open.**

***Index Terms*—speaker embeddings, speaker diarization, speaker change detection, code-switching, Nepali, cross-lingual speaker verification, pre-specified analysis**

I. Introduction

Speaker diarization relies on speaker embeddings and on scores that compare adjacent windows. If a language switch by one speaker (SS-DL: *same speaker, different language*) moves the embedding, a diarization system could report a speaker change that is not there. We test the first link of that chain. Is same-speaker embedding similarity lower, and a two-window change score higher, at natural Nepali↔English switches than at matched same-language boundaries of the same speaker? And do mitigations meant to remove language information from the embedding reduce that excess?

Cross-lingual speaker verification already shows that embeddings rate the same speaker as less similar across languages [1], [2], and LASE [3] reports the same for synthetic Indic voices. Diarization studies of code-mixed speech report corpus-level error [4]. What is missing is a measurement *at individual, natural switch points* in real speech, against matched same-speaker controls. We provide one for Nepali–English. LASE also found that the cross-script gap is smaller for Indian-accented voices, so for speakers whose first language is Nepali the effect may be small; we fixed the analysis plan before the confirmatory run so that a small or null result could still be reported.

In short, between language runs of at least a second, same-speaker similarity is slightly but consistently lower across a switch, and the change score slightly higher. Over all switches, most of which border a run shorter than half a second, there is no positive ranking separation, and neither testable mitigation demonstrably removed the excess we found. The planned tests on full diarization pipelines and on real cross-lingual speaker turns (DS-DL: *different speaker, different language*) needed conversation data we could not obtain, so this is a diagnostic study of embeddings and change scores, with a post hoc pipeline check.

Contributions.

- A boundary-level protocol that compares speaker-embedding similarity and a two-window change score at same-speaker language switches with matched same-language controls, adjusted for measured acoustic covariates.
- Pre-specified results on real Nepali–English speech for three speaker embeddings (§V), with a check of direction on Hindi– and Bengali–English (§VI).
- A pre-specified test of two mitigations meant to remove language information (§VII) and a post hoc check with two off-the-shelf diarization pipelines (§VIII).

Code, the analysis plan and results are at github.com/ishangautam7/nepali-cs-diarization.

II. Related Work

Code-switching. Linguistic work distinguishes the insertion of single words or constituents into a matrix language from alternation between longer stretches of two languages [5], [6]. Doğruöz et al. [7] survey code-switching for language technologies, and the code-mixing index (CMI) measures how mixed a text is [8]. Most of our switch points are insertion-like (§IX).

Cross-lingual speaker representations. Speaker embeddings underestimate same-speaker similarity across languages, and probing shows that they also encode spoken content and

channel [9]. Thienpondt et al. [1] correct this score shift with language information. Later work removes language with subspace normalisation [10] or gradient-reversal adversarial training [2], [3]. Hosseini-Kivanani [10] and Li et al. [2] evaluate verification only; Li et al. train on real multilingual speech and found that added TTS speech did not improve verification. LASE [3] trains on TTS-generated cross-script pairs and adds a synthetic diarization benchmark of 50 conversations built from 2–4 TTS voices, one of which may switch language, scored by clustering agreement with segments and speaker count given. It also reports that the cross-script gap of off-the-shelf encoders is smaller for synthetic Indian-accented voices (0.006 vs. 0.082 for WavLM-SV; 0.044 vs. 0.105 for ECAPA-TDNN).

Diarization. Modern systems combine a powerset segmentation model [11] with embedding clustering: DiariZen pairs WavLM-based segmentation with WeSpeaker embeddings and VBx-style clustering [12], [13], and DiaPer uses perceiver attractors [14]. The DISPLACE challenges [4], [15] report diarization error on multilingual Indian conversations, DISPLACE-M extends this to code-switched medical speech [16], [17], and the closest Nepali work evaluates multilingual diarizers on Nepali–Hindi speech [18]. These studies report corpus-level error; none of them, nor LASE's synthetic benchmark, measures embedding similarity or a change score at individual natural switch points. Mishra and Prasanna [19] detect spoken-language change by analogy with speaker-change detection, which suggests that the two are acoustically entangled; our acoustic covariates and pause controls address part of this.

III. Data

A. D1: Nepali–English lectures (main corpus)

We use a Nepali–English code-switched corpus of STEM lectures [20]: 20,805 clips (58.4 h, 3–14 s each) from 161 videos across 20 YouTube channels, at 16 kHz. Its transcripts are mostly machine-generated (777 clips are human-verified), so language tags and switch points inherit transcription errors. The dataset's speaker label is the *channel*, not a verified person, so we audit speaker purity (§IV-C). We use only the real (non-augmented) clips. The 3 validation channels are the pilot, used to fix every threshold; the confirmatory analysis runs once on the other 17 channels (15 training, 2 test).

B. D4: MUCS 2021 Hindi/Bengali–English (direction check)

To check the direction of the effect in other language pairs, we use the MUCS 2021 code-switched subsets [21]: Hindi–English (52,825 train and 3,136 test clips) and Bengali–English (26,606 and 4,275), cut from narrated computer-science tutorials. The MUCS authors had no speaker information and assume that each tutorial has a different speaker [21], so a MUCS "speaker" is a tutorial and one narrator may appear under many IDs. MUCS often writes English words in native script (e.g. "office" in Devanagari), so we tag language with a per-script lexicon of 763 loans and 67 names for Hindi and 393 and 21 for Bengali, labelled by an LLM and spot-checked as text. This gives 4,267 (Hindi) and 2,769 (Bengali) switch points with matched controls for H1, and 5,071 and 3,510 with cell controls in the primary detection cell.

C. Data we could not obtain

The planned tests of full diarization pipelines need long-form recordings, and the DS-DL tests need conversations with references for both speaker and language. D1's source videos could not be downloaded, access to DISPLACE 2024 [15] was not granted during the study, and DISPLACE-M [16] has no language references. These tests were therefore not run: the full-pipeline variants of the change-score test (H2), whether the mitigations add missed cross-language speaker turns (H5), and their ordering (H6). Apart from the artificial speaker changes of §VIII, no result here involves a speaker change.

IV. Method

A. Switch points, controls and matching

We romanise each word with uroman [22], force-align each clip with torchaudio's MMS_FA bundle (MMS-300M [23]), using Whisper [24] word timestamps as a cross-check, tag each word's language from its script with a loan/name lexicon, and detect pauses with Silero VAD [25]. An SS-DL switch point is the boundary between two adjacent words of the same speaker in different languages, with both words aligned and no neutral or loan word nearby. Each point is assigned to a run-length stratum (< 0.5, 0.5–1, ≥ 1 s) by the shorter adjacent language run.

Two analysis families use the points; a third test (H3, §IV-E) compares cross-language with mono-language same-speaker trials. The *embedding* family (H1) compares each switch point between runs of at least 1 s with up to three same-language control boundaries from the same clip or video, matched exactly on pause bin, position bin and context length; its 1 s windows exclude any pause. The *detection* family (H2-P0, for the change score P0 below) compares switch points with all same-speaker controls in the same video $\times$ pause-bin $\times$ position-bin cell, standardised per cell; its 1 s windows meet at the middle of any pause.

Primary detection cell. The primary cell, `p0L1`, keeps switch points where both language runs last at least 1 s, with a clean ± 1 s field: on the confirmatory channels, 3,962 switch points with cell controls, against 31,844 controls. The pooled cell holds all 120,630 switch points. The initial study design, committed before the pilot, made the pooled cell primary. On the pilot channels its AUC was 0.486–0.494, while switch points between longer runs fired more often than controls. We then made `p0L1` primary, because a detector with 1 s windows cannot register a sub-second insertion as a two-sided change; this also gives P0 the run-length condition the initial design already fixed for H1. This change was made after seeing pilot outcomes. The formal analysis plan was written after it, names `p0L1` as primary, and was frozen with the configuration before the one confirmatory run; the confirmatory channels are disjoint from the pilot. Because the primary cell shows

a positive result where the pooled cell shows no positive separation, we report both.

B. Systems

We test three speaker embeddings: E1 ECAPA-TDNN [26], E2 the WeSpeaker ResNet34-LM embedding [27] of pyannote Community-1, and E3 WavLM-Large + ECAPA [28]. The change score P0 is $d(t) = 1 - \cos$ between the 1 s windows before and after each position on a dense grid. Within ± 0.5 s of a boundary it keeps the highest peak and *fires* when that peak exceeds an operating point τ, set on the pilot so that 10% of controls fire. ΔFAR is how much more often P0 fires at switch points than at their cell controls. Every fire at a same-speaker boundary is a false alarm, so ΔFAR is a change-score false-alarm excess, not a diarization error rate. In the primary cell P0 compares windows much like H1's, so its AUC largely restates H1's similarity drop as a ranking; what P0 adds is peak picking and an operating point.

To check that switch points carry a language change, we scored each boundary on the pilot channels by the Jensen–Shannon divergence between language-ID posteriors [29] of the 1 s windows on either side. It ranked switch points above controls with AUC 0.58 overall and 0.72 where the shorter run lasts at least 1 s, but it also ranks pause-bearing controls above other controls (AUC 0.56), so it is only a weak check.

C. Speaker-purity audit

Because a speaker is a channel, we audit each video with two embeddings not under test (TitaNet-L and CAM++), comparing each clip only with same-video clips of similar language mix (by CMI tertile [8]), so that English-heavy clips are not flagged for their language. The audit flagged 688 of the 20,444 clips assessable by both embeddings (3.4%), at the same rate in every language-mix stratum, and split off 34 of 161 videos, most of them separate recording sessions of one teacher. Apart from one pilot-channel video that may share a teacher with a confirmatory channel, the closest pair of videos from two channels has a cosine of 0.65 (TitaNet-L) and 0.66 (CAM++), below the 0.77–0.95 typical of two videos of one channel. In a blind listening check of 117 pairs by the author, the main voice of a flagged clip was judged to differ from its video's usual speaker in 51% of usable pairs (95% CI 41–62%), against 8% for unflagged control pairs (1 of 12; 95% CI 1–35%). With flagged clips dropped and split-off videos as separate clusters, every primary result keeps its direction and meets its smallest effect of interest (H1 d_z 0.25/0.21/0.26, H2-P0 AUC 0.58/0.56/0.59, H3 $-0.24/-0.20/-0.25$ SD for E1/E2/E3).

D. Mitigations

The plan specified three mitigations and three baselines. M1 projects the language direction out of the embedding [10]. M2 gates the change score with language ID [29]. M3 fine-tunes E3 on real speech only, with a gradient-reversal language loss [30], AAM-softmax [31], a supervised-contrastive loss and distillation towards E3, adding a Nepali speaker pool from IndicVoices [32], with five seeds per fold. The baselines are B0, the unmodified system with its operating point recalibrated; B1, language-wise mean centring; and B2, a language-aware score offset [1]. The test (H4) uses P0 with E3 on the 15 training channels in 5 folds of 3: every component is fitted on the other 12 channels and every τ is set on the pilot. For each channel, $\delta_M = \Delta\mathrm{FAR}_M - \Delta\mathrm{FAR}_{B0}$ is tested by an exact one-sided sign-flip with Holm correction over M1–M3; support also needs a reduction of at least 25% of B0's excess. M2 could not be tested: its gate weight was to be tuned on conversations, because on D1, which has no speaker changes, any gate lowers false alarms at no visible cost. For the same reason, a lower excess alone is not an improvement.

E. Statistics

The unit of analysis is the channel. We use exact sign-flip permutation tests over the 17 confirmatory channels (2^{17} sign patterns), BCa bootstrap intervals at $1 - 0.05/m$, and Holm correction within each hypothesis family ($m = 3$ for H1 and H3; $m = 5$ for H2, whose untested pipeline variants count at $p = 1$). For H1, d_z is the mean over channels of the mean cosine drop (control minus switch), divided by the pooled within-channel SD of the drops. For H2-P0, a channel's AUC is the probability that a switch point's $d(t)$ exceeds a same-cell control's. H3 is the mean target-score difference (cross-language minus mono-language) on lag-matched same-speaker trials, in units of target-score SD, with the two languages weighted equally. The smallest effects of interest are $d_z = 0.2$, AUC $= 0.55$ and 0.2 SD; a result is *supported* if it is significant with a point estimate at least this large. The sample was fixed by the data; with pilot variance components, the planned power at these effects was 0.68–1.00 for H1 and 0.70–0.76 for H2-P0.

An excess is *covariate-robust* if the primary test, a channel-level test adjusted for acoustic covariates (left–right changes in F0, energy, speaking rate, spectral tilt and SNR), a contrast against same-language phrase boundaries (pause > 150 ms) and H3 are all significant; otherwise it is *switch-associated*. Reverberation, prosodic phrasing and syntactic position were not measured, so this grade shows an association that survives adjustment, not that language causes it.

V. Results on D1

All nine primary tests are significant (Table I), and every channel shows the predicted direction for H1 and H2-P0 with all three embeddings, and for H3 in 16 or 17 of 17. The tests are not independent: H1 and H2-P0 compare similar windows. For E2, H1 and H3 are significant with point estimates just below the smallest effect of interest, which their intervals still include.

The excess depends on run length. Over all switch points there is no positive ranking separation (pooled AUC 0.485–0.493), and at τ the pooled ΔFAR is negative ($-1.59/-1.67/-1.51$ pp for E1/E2/E3; secondary, no test). The pooled AUC rises from ≈ 0.48 when the shorter run is under 0.5 s to ≈ 0.59 when it is at least 1 s (exploratory), consistent with a

TABLE I

D1 PRIMARY OUTCOMES ON 17 CHANNELS. HOLM-ADJUSTED SIGN-FLIP TESTS; UNADJUSTED ESTIMATES WITH BCA CIS AT $1 - 0.05/m$ (98.3% FOR H1 AND H3, 99% FOR H2-P0). †SIGNIFICANT, POINT ESTIMATE BELOW THE SMALLEST EFFECT OF INTEREST. H1 USES 3,695 MATCHED SETS; H2-P0 USES THE PRIMARY CELL (§IV).

H	Model	Estimate (CI)	Outcome
H1	E1	d_z 0.271 (0.203–0.348)	supported
H1	E2	d_z 0.194 (0.135–0.266)	significant†
H1	E3	d_z 0.254 (0.198–0.305)	supported
H2-P0	E1	AUC 0.584 (0.559–0.618)	supported
H2-P0	E2	AUC 0.573 (0.547–0.604)	supported
H2-P0	E3	AUC 0.582 (0.564–0.601)	supported
H3	E1	−0.215 SD (−0.240–−0.175)	supported
H3	E2	−0.173 SD (−0.204–−0.124)	significant†
H3	E3	−0.203 SD (−0.239–−0.152)	supported

TABLE II

CHANGE-SCORE DETECTION. P0L1 IS THE PRIMARY CELL; THE POOLED CELL (ALL SWITCH POINTS) AND ΔFAR ARE SECONDARY. ΔFAR IS THE EQUAL-WEIGHT MEAN OVER CHANNELS (FIG. 1), WITH ITS ONE-SIDED SIGN-FLIP p.

Model	p0L1 AUC	pooled AUC	ΔFAR at τ (p)
E1	0.584	0.488	+2.46 pp (0.046)
E2	0.573	0.485	+2.95 pp (0.023)
E3	0.582	0.493	+2.73 pp (0.003)

1 s detector being unable to register short insertions. Window length behaves similarly for H1: with 0.5 s windows d_z falls to 0.13/0.10/0.12 (E1/E2/E3), and with 1.5 s windows it rises to 0.34/0.24/0.36, though longer windows also admit only longer runs.

Size in practical terms. On H3's trials the shift costs 0.6–0.9 pp of equal error rate (EER) (12.9/9.5/11.4% cross-language against 12.2/8.9/10.5% mono-language). At τ, P0 fires at 13–17% of primary-cell switch points but at 22–25% of same-language boundaries with a pause over 150 ms (channel means, not cell-standardised), so for this detector a pause is a larger source of within-speaker false alarms than a language switch. The operating-point excess is also uneven across channels (Fig. 1): ΔFAR is positive in 9, 12 and 13 of 17 channels, and weighting switch points rather than channels equally gives +0.78, +2.00 and +1.17 pp, because the largest channel has a negative ΔFAR and the largest values come from small channels.

Acoustic controls. H1 is covariate-robust for all three embeddings. The H2-P0 excess is covariate-robust for E2 ($p = .031$) and E3 ($p = .008$) and switch-associated for E1 ($p = .079$). A descriptive mixed model agrees for H1 (switch effect $-0.030/-0.024/-0.031$ in cosine, $|z| \geq 7.3$, almost unchanged without covariates) but finds the H2-P0 excess weak (one-sided p .081/.026/.124).

Robustness. Every outcome category is unchanged when low-confidence alignments are dropped or the 2 test channels are left out, and no H2-P0 or H3 category changes when switch points are re-extracted under three alternative rules (two tagging changes, and keeping low-alignment clips); only

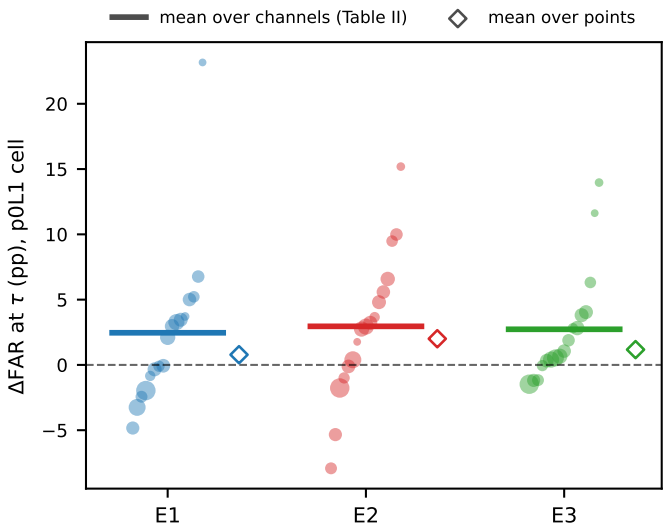


Fig. 1. Change-score excess at the operating point, per confirmatory channel (primary cell). Each dot is one channel's ΔFAR, the fire rate at its switch points minus that at its matched controls; dot area grows with the channel's number of switch points (13 to 584). Bars mark the equal-weight mean over channels (Table II); diamonds mark the mean over switch points.

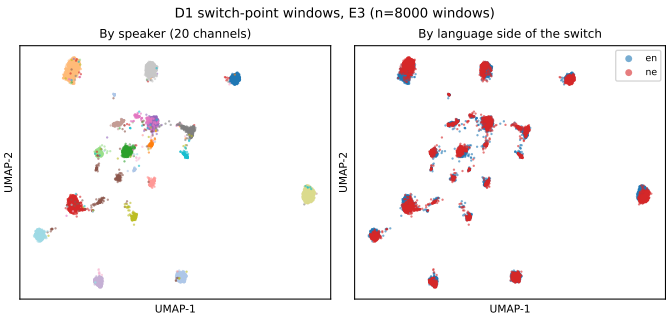


Fig. 2. UMAP [33] of 8000 D1 switch-point windows (E3). *Left:* coloured by channel; a linear probe identifies the channel with balanced accuracy 0.95 (cross-session probe over six channels at a time, chance 0.17). *Right:* coloured by language; the languages interleave within each channel with a small offset, and a linear probe still recovers language with balanced accuracy 0.76 (chance 0.5). Illustrative only; the probe accuracies come from the mitigation folds, not from the plotted sample.

E2's H1, which sits at its smallest effect throughout (d_z 0.17–0.21), moves between grades, when the low-alignment clips are kept. Leaving out any one channel (post hoc) reverses no estimate and leaves every primary test significant; the least stable quantity is E1's ΔFAR, which falls to +1.17 pp ($p = .09$) without its smallest channel (13 switch points).

Listening check. The author listened blind to 352 items. No language change was audible at 6.3% of primary-cell switch points (95% CI 2.0–11.6%), none had a different speaker or overlap, and no control had a language change. Correcting ΔFAR for this contamination raises it by about 7%. Fig. 2 illustrates the embedding geometry behind H1 and H3.

TABLE III
DIRECTION CHECK ON D4 WITH TUTORIALS AS CLUSTERS. ALL HOLM ONE-SIDED $p < .05$. "SUPPORTED/3" COUNTS EMBEDDINGS AT OR ABOVE D1'S SMALLEST EFFECT. CIS ARE 99.58% BCA.

Pair	H	Estimate E1/E2/E3 (CI range)	supported/3
hi-en	H1	d_z 0.24 / 0.18 / 0.26 (.13–.31)	2/3
hi-en	H2-P0	AUC 0.56 / 0.56 / 0.58 (.54–.60)	3/3
bn-en	H1	d_z 0.16 / 0.12 / 0.15 (.05–.24)	0/3
bn-en	H2-P0	AUC 0.53 / 0.52 / 0.54 (.50–.57)	0/3

VI. DIRECTION CHECK ON HINDI– AND BENGALI–ENGLISH

We repeat H1 and H2-P0 on D4, specified and frozen before any D4 outcome was computed, with MUCS speaker IDs (tutorials) as clusters. A tutorial enters a test if it has at least 5 matched sets (H1) or 5 primary-cell switch points (H2-P0): 364 and 394 of 550 Hindi IDs, and 235 and 254 of 307 Bengali IDs. With hundreds of clusters, the sign-flip draws 2^{17} random sign patterns. All 12 tests are significant in the predicted direction (Holm one-sided $p \leq 2.6 \times 10^{-3}$; Table III). In Hindi–English the estimates are comparable to D1; in Bengali–English every point estimate falls below D1's smallest effect, so there the direction matches, but the effect is not shown to reach D1's practical size.

These p-values assume that tutorials are independent speakers; if one narrator recorded many tutorials, they are too small. The planned check groups recordings into voice clusters, but its threshold, calibrated on D1, merged MUCS into only 6 Hindi and 3 Bengali clusters. With so few clusters no Holm-adjusted p can fall below .05 (the smallest attainable is ≈ 0.19), so no test can reach significance; this reflects a lack of power, not evidence against the effect, and the check is uninformative. Robustness to speaker dependence is therefore untested on D4, which supports the direction of the effect, not its significance under realistic speaker dependence. An LLM-labelled text audit of 200 native-tagged tokens per pair put residual English at 0.5% for Hindi and 1.0% for Bengali; re-reading the Bengali items found four missed English words, which raises it to 3.0%, probably the fairer figure.

VII. MITIGATIONS

On the 15 scored channels B0's excess is $+2.13$ pp, so support needs a reduction of at least 0.533 pp. Neither mitigation is supported (Table IV): M1 lowers the excess by 0.40 pp (18.9%) and M3 by 0.53 pp (24.9%), but both intervals include no reduction. The language-aware offset B2, itself not significant ($p = .10$), lowers it by about as much (0.59 pp), and mean centring B1 does not lower it.

Linear probes explain part of this null. M1 removes much of the decodable language information (language balanced accuracy 0.76 for E3, 0.62 after M1, with speaker identification unchanged at 0.95), yet the excess barely falls. M3, despite its adversarial loss, makes language *more* decodable (0.80–0.82 over seeds) while improving verification (mono-language EER 0.118 for E3, 0.062–0.071 for M3). M3 therefore did not become language-invariant, and H4 does not test whether a language-invariant embedding would help. In single-seed ablations, removing the adversarial loss kept M3's small reduction, while removing distillation reversed it, so the reduction does not seem to come from the adversarial loss. On the 2 held-out test channels (descriptive), B0's excess is $+7.22$ pp and neither mitigation lowers it consistently.

TABLE IV
MITIGATIONS (H4): CHANGE IN ΔFAR AGAINST B0 ($+2.13$ PP) OVER 15 CHANNELS, M3 AVERAGED OVER 5 SEEDS PER FOLD. ONE-SIDED SIGN-FLIP, HOLM OVER M1–M3; BCA CIS AT 98.3%.

System	Change (pp)	CI (pp)	Holm p	Outcome
M1	-0.40 (18.9%)	-0.98 to $+0.17$	0.20	inconclusive
M2	—	—	1	not tested
M3	-0.53 (24.9%)	-3.89 to $+0.74$	0.70	inconclusive

TABLE V
POST HOC PIPELINE CHECK: SHARE OF 300 BOUNDARIES AT WHICH THE PIPELINE REPORTS A SPEAKER CHANGE. FOR ONE SPEAKER THIS IS A FALSE CHANGE; FOR TWO, ONE MINUS IT IS THE MISS RATE. JOINS ARE ARTIFICIAL.

Boundary	Speakers	Community-1	DiariZen
Natural switch	one	1.3%	0.3%
Natural same-language control	one	0.3%	0.3%
Join, one video	one	1.7%	3.3%
Join, two channels, same language	two	60.3%	94.0%
Join, two channels, language change	two	65.3%	96.0%

VIII. EXPLORATORY FULL-PIPELINE CHECK

After all planned analyses, we ran two off-the-shelf diarization pipelines with default settings and no given speaker count: pyannote Community-1, with its VBx clustering [13], and DiariZen [12]. On the 17 confirmatory channels we drew 300 primary-cell switch points and one same-cell control for each, and ran each pipeline on the whole clip. Because D1 has no speaker changes, we also built artificial ones: 300 joins of clips from two channels with the same language at the join, and 300 with a language change; 300 joins of consecutive clips from one video serve as same-speaker joins. A join also carries a recording cut, so it tests whether a speaker change is found, not natural turn-taking. A boundary counts as a reported change if the output switches speaker within ± 0.5 s.

Neither pipeline reported a speaker change at natural switches significantly more often than at controls (mean per-channel differences $+1.2$ and -0.4 pp; two-sided $p = .31$ and 1.0; Table V), so the change-score excess did not appear as a detectable excess of false speaker changes in these short clips. Community-1 missed 35–40% of the artificial speaker changes, so its null says little; DiariZen found 94–96% of them and gives the more informative null.

IX. DISCUSSION

On real Nepali–English speech, same-speaker embedding similarity is lower across a language switch (H1), and cross-language same-speaker scores are lower (H3), confirming at

natural switch points the shift known from verification [1], [3]. The shift survives the measured acoustic and pause controls, and it is small: standardised effects of 0.17–0.27, or a cosine drop of 0.037/0.029/0.034 (E1/E2/E3). For ECAPA-TDNN this is close to LASE's off-the-shelf gap for synthetic Indian-accented voices (0.044) and well below that for Western-accented voices (0.105), although LASE's gaps are utterance-level and ours come from 1 s windows.

The change score tells a narrower story. It is higher at switches only where both language runs last at least a second. Most switches in these lectures are insertion-like [6]: 79% of the 120,630 pooled switch points border a run shorter than 0.5 s, too short for a 1 s detector to resolve, and over all switches there is no positive ranking separation. The two off-the-shelf pipelines rarely reported a false change at switches. Neither mitigation demonstrably removed the excess, and neither did better than a simple score offset. In code-switched lectures, then, language switches by one speaker look like a minor source of false speaker changes. Whether that holds in conversation, and whether a mitigation would cost missed real speaker changes, needs conversations with speaker and language references.

Limitations. No planned pipeline or DS-DL test could be run, so the results mainly concern embeddings and change scores. The primary detection cell was chosen after the pilot. The design is observational. D1's speakers are YouTube channels and D4's are tutorials, neither verified, and the channel-level tests assume they are independent. The listening checks had one rater, the author, who knew the hypotheses; a subset for two independent bilingual raters is prepared but not yet labelled, so the evidence is preliminary. D1's transcripts are mostly machine-generated, and the planned measurement of tag noise on its 777 human-verified clips was not run; the D4 language labels were checked only as text. Both corpora are instructional monologues without speaker metadata such as gender, and switching in conversation may differ in frequency, run length and prosody. With 15 channels, H4 could not separate a 25% reduction from none.

X. Conclusion

Within one speaker and between language runs of at least a second, three speaker embeddings rate the same voice as slightly less similar across a natural Nepali–English switch, and a two-window change score is slightly higher there. Over all switches there is no positive ranking separation, two mitigations reduced the excess only inconclusively, and two off-the-shelf pipelines rarely reported a false speaker change at switches. Testing the effect in conversational diarization needs long-form conversations with speaker and language references.

XI. Ethics, Data Availability and Use of AI

The study uses only publicly released recordings: D1 (CC BY-NC 4.0) [20], MUCS 2021 (CC BY-SA 4.0) and the Nepali part of IndicVoices (CC BY 4.0). It makes no new recordings, recruits no participants and redistributes no audio, so no ethics-board review was sought. Speaker embeddings are biometric data, so only aggregate statistics are reported. Code, the analysis plan and all results are available at the repository above. The plan was written and committed to the project's git repository before the confirmatory analyses, and each confirmatory analysis was frozen as a tagged commit, recording a hash of every code and input file, before it was run. We call the plan pre-specified rather than pre-registered because it was not deposited with an external registry; the full commit history is available on request. Post hoc analyses are labelled as such.

The author designed the study, wrote the analysis plan, made the analysis decisions, did the listening checks and interpreted the results. AI assistants (Anthropic Claude models) helped write and run the analysis code and edit the text. LLM-based tools labelled the D4 lexicon and its audit sample, as text only. The author checked their output and is responsible for all content.

Conflicts of Interest

The author declares no conflicts of interest.

References

[1] J. Thienpondt, B. Desplanques, and K. Demuynck, "Tackling the score shift in cross-lingual speaker verification by exploiting language information," in *Proc. IEEE ICASSP*, 2022, arXiv:2110.09150.

[2] Z. Li, X. Miao, J. Liu, and M. Li, "Language-invariant multilingual speaker verification for the TidyVoice 2026 challenge," 2026, arXiv:2603.08092.

[3] V. P. T. Menta, "LASE: Language-adversarial speaker encoding for Indic cross-script identity preservation," 2026, arXiv:2605.00777.

[4] S. Baghel *et al.*, "DISPLACE challenge: DIarization of SPeaker and LAnguage in conversational environments," in *Proc. Interspeech*, 2023, arXiv:2303.00830.

[5] S. Poplack, "Sometimes I'll start a sentence in Spanish Y TERMINO EN ESPAÑOL: Toward a typology of code-switching," *Linguistics*, vol. 18, no. 7–8, pp. 581–618, 1980.

[6] P. Muysken, *Bilingual Speech: A Typology of Code-Mixing*. Cambridge, UK: Cambridge University Press, 2000.

[7] A. S. Doğruöz, S. Sitaram, B. E. Bullock, and A. J. Toribio, "A survey of code-switching: Linguistic and social perspectives for language technologies," in *Proc. ACL-IJCNLP*, 2021, pp. 1654–1666.

[8] B. Gambäck and A. Das, "Comparing the level of code-switching in corpora," in *Proc. LREC*, 2016, pp. 1850–1855.

[9] D. Raj, D. Snyder, D. Povey, and S. Khudanpur, "Probing the information encoded in X-vectors," in *Proc. IEEE ASRU*, 2019, pp. 726–733, arXiv:1909.06351.

[10] N. Hosseini-Kivanani, "Simple language normalization wins: Cross-lingual speaker verification for the TidyVoice 2026 challenge," 2026, arXiv:2607.22923.

[11] A. Plaquet and H. Bredin, "Powerset multi-class cross entropy loss for neural speaker diarization," in *Proc. Interspeech*, 2023.

[12] J. Han, F. Landini, J. Rohdin, A. Silnova, M. Diez, and L. Burget, "Leveraging self-supervised learning for speaker diarization," in *Proc. IEEE ICASSP*, 2025, arXiv:2409.09408.

[13] F. Landini, J. Profant, M. Diez, and L. Burget, "Bayesian HMM clustering of x-vector sequences (VBx) in speaker diarization: Theory, implementation and analysis on standard tasks," *Computer Speech & Language*, vol. 71, p. 101254, 2022, arXiv:2012.14952.

[14] F. Landini, M. Diez, T. Stafylakis, and L. Burget, "DiaPer: End-to-end neural diarization with Perceiver-based attractors," *IEEE/ACM Trans. Audio, Speech, Language Process.*, 2024, arXiv:2312.04324.

[15] S. B. Kalluri *et al.*, "The second DISPLACE challenge: DIarization of SPeaker and LAnguage in conversational environments," in *Proc. Interspeech*, 2024, arXiv:2406.09494.

[16] Dhanya E, A. Meena, M. Nanivadekar *et al.*, "Benchmarking speech systems for frontline health conversations: The DISPLACE-M challenge," 2026, arXiv:2603.02813.

[17] S. Barahona, L. Herrera-Alarcón, J. I. Alvarez-Trejos, and A. Lozano-Diez, "Adapting speaker diarization to code-switched medical conversations: AUDIAS-UAM at the DISPLACE-M challenge," in *Proc. Odyssey*, 2026, pp. 75–82.

[18] S. Neupane, S. Pokhrel, S. Pyakurel, and B. Joshi, "Neural speaker diarization via multilingual training: Evaluation on low-resource Nepali-Hindi speech," 2026, arXiv:2606.26144.

[19] J. Mishra and S. R. M. Prasanna, "Spoken language change detection inspired by speaker change detection," 2023, arXiv:2302.05265.

[20] S. Dahal, "Nepali–English code-switched ASR (dataset, v2)," 2026, Hugging Face dataset saileshbro/nepali-cs-asr, revision dd00ca41. [Online]. Available: https://huggingface.co/datasets/saileshbro/nepali-cs-asr

[21] A. Diwan *et al.*, "MUCS 2021: Multilingual and code-switching ASR challenges for low resource Indian languages," in *Proc. Interspeech*, 2021, arXiv:2104.00235.

[22] U. Hermjakob, J. May, and K. Knight, "Out-of-the-box universal romanization tool uroman," in *Proc. ACL, System Demonstrations*, 2018, pp. 13–18.

[23] V. Pratap *et al.*, "Scaling speech technology to 1,000+ languages," *J. Mach. Learn. Res.*, vol. 25, no. 97, pp. 1–52, 2024, arXiv:2305.13516.

[24] A. Radford *et al.*, "Robust speech recognition via large-scale weak supervision," in *Proc. ICML*, 2023, arXiv:2212.04356.

[25] Silero Team, "Silero VAD: Pre-trained enterprise-grade voice activity detector," https://github.com/snakers4/silero-vad, 2021.

[26] B. Desplanques, J. Thienpondt, and K. Demuynck, "ECAPA-TDNN: Emphasized channel attention, propagation and aggregation in TDNN based speaker verification," in *Proc. Interspeech*, 2020, arXiv:2005.07143.

[27] H. Wang *et al.*, "Wespeaker: A research and production oriented speaker embedding learning toolkit," in *Proc. IEEE ICASSP*, 2023, arXiv:2210.17016.

[28] S. Chen *et al.*, "WavLM: Large-scale self-supervised pre-training for full stack speech processing," *IEEE J. Sel. Topics Signal Process.*, vol. 16, no. 6, pp. 1505–1518, 2022, arXiv:2110.13900.

[29] J. Valk and T. Alumäe, "VoxLingua107: a dataset for spoken language recognition," in *Proc. IEEE SLT*, 2021, arXiv:2011.12998.

[30] Y. Ganin and V. Lempitsky, "Unsupervised domain adaptation by backpropagation," in *Proc. ICML*, 2015.

[31] J. Deng *et al.*, "ArcFace: Additive angular margin loss for deep face recognition," in *Proc. IEEE CVPR*, 2019.

[32] T. Javed *et al.*, "IndicVoices: Towards building an inclusive multilingual speech dataset for Indian languages," in *Findings of ACL*, 2024, arXiv:2403.01926.

[33] L. McInnes, J. Healy, and J. Melville, "UMAP: Uniform manifold approximation and projection for dimension reduction," 2018, arXiv:1802.03426.